

ThoughtWorks®

Workshop

ELASTICSEARCH

Felipe Dornelas

AGENDA

☐ *Part 1*

- ☐ Introduction
- ☐ Document Store
- ☐ Search Examples
- ☐ Data Resiliency
- ☐ Comparison with Solr

☐ *Part 2*

- ☐ Search
- ☐ Analytics

AGENDA

☐ *Part 3*

- ☐ Inverted Index
- ☐ Analyzers
- ☐ Mapping
- ☐ Proximity Matching
- ☐ Fuzzy Matching

☐ *Part 4*

- ☐ Inside a Cluster
- ☐ Data Modeling

→ [github.com/felipead/
elasticsearch-workshop](https://github.com/felipead/elasticsearch-workshop)

PRE-REQUISITES

- ☐ Vagrant
- ☐ VirtualBox
- ☐ Git

ENVIRONMENT SETUP

- ❑ `git clone https://github.com/felipead/elasticsearch-workshop.git`
- ❑ `vagrant up`
- ❑ `vagrant ssh`
- ❑ `cd /vagrant`

VERIFY EVERYTHING IS WORKING

❑ `curl http://localhost:9200`

ThoughtWorks®

PART 1

Core concepts

ThoughtWorks®

1-1 INTRODUCTION

You know, for search

WHAT IS ELASTICSEARCH?

*A real-time distributed **search** and
analytics engine*

IT CAN BE USED FOR

- ❑ *Full-text* search
- ❑ *Structured* search
- ❑ Real-time *analytics*
- ❑ ...or any combination of the above

FEATURES

- ❑ Distributed *document store*:
 - ❑ *RESTful API*
 - ❑ Automatic scale
 - ❑ Plug & PlayTM

FEATURES

- ❑ Handles the *human language*:
 - ❑ Score results by *relevance*
 - ❑ *Synonyms*
 - ❑ *Typos* and *misspellings*
 - ❑ Internationalization

FEATURES

- ❑ Powerful *analytics*:
 - ❑ Comprehensive *aggregations*
 - ❑ Geolocations
 - ❑ Can be combined with *search*
 - ❑ *Real-time* (no batch-processing)

FEATURES

- ❑ *Free and open source*
- ❑ Community support
- ❑ Backed by Elastic

MOTIVATION

*Most databases are inept at extracting **knowledge** from your data*

SQL DATABASES

SQL = Structured Query Language

SQL DATABASES

- ❑ Can only filter by *exact values*
- ❑ *Unable* to perform *full-text* search
- ❑ Queries can be *complex* and *inefficient*
- ❑ Often requires *big-batch* processing

APACHE LUCENE

- ❑ Arguably, the *best* search engine
- ❑ High performance
- ❑ Near real-time indexing
- ❑ Open source


APACHE LUCENE

- But...

- It's just a *Java Library*

- Hard to use

ELASTICSEARCH



- ❑ Document Store
- ❑ Distributed
- ❑ Scalable
- ❑ Real Time
- ❑ Analytics
- ❑ RESTful API
- ❑ *Easy to Use*

DOCUMENT ORIENTED

- ❑ *Documents* instead of rows / columns
- ❑ Every field is *indexed* and *searchable*
- ❑ Serialized to *JSON*
- ❑ *Schemaless*

WHO USES

- ❑ GitHub
- ❑ Wikipedia
- ❑ Stack Overflow
- ❑ The Guardian

TALKING TO ELASTICSEARCH

□ *Java API*

- Port 9300
- Native transport protocol
- Node client (*joins the cluster*)
- Transport client (*doesn't join the cluster*)

TALKING TO ELASTICSEARCH

- *RESTful API*
 - Port 9200
 - JSON over HTTP

TALKING TO ELASTICSEARCH

*We will only cover the **RESTful** API*

USING CURL

`curl -X <VERB> <URL> -d <BODY>`

or

`curl -X <VERB> <URL> -d @<FILE>`

THE EMPTY QUERY

```
curl -X GET
```

```
-d @part-1/empty-query.json
```

```
localhost:9200/_count?pretty
```

REQUEST



```
{  
  "query": {  
    "match_all": {}  
  }  
}
```

RESPONSE



```
{  
  "count": 0,  
  "_shards": {  
    "total": 0,  
    "successful": 0,  
    "failed": 0  
  }  
}
```

1-2 DOCUMENT STORE

THE PROBLEM WITH RELATIONAL DATABASES

- ❑ Stores data in *columns* and *rows*
- ❑ Equivalent of using a *spreadsheet*
- ❑ *Inflexible* storage medium
- ❑ Not suitable for *rich objects*

DOCUMENTS

```
{  
  "name": "John Smith",  
  "age": 42,  
  "confirmed": true,  
  "join_date": "2015-06-01",  
  "home": {"lat": 51.5, "lon": 0.1},  
  "accounts": [  
    {"type": "facebook", "id": "johnsmith"},  
    {"type": "twitter", "id": "johnsmith"}  
  ]  
}
```

DOCUMENT METADATA

- *Index* - Where the document lives
- *Type* - Class of object that the document represents
- *Id* - Unique identifier for the document

DOCUMENT METADATA

*Relational
DB*

Databases

Tables

Rows

Columns

Elasticsearch

Indices

Types

Documents

Fields

RESTFUL API

[VERB] /{index}/{type}/{id}?pretty

GET | POST | PUT | DELETE | HEAD

RESTFUL API

- ❑ *JSON*-only
- ❑ Adding *pretty* to the query-string
parameters pretty-prints the response

INDEXING A DOCUMENT WITH YOUR OWN ID

PUT `/`{index}`/`{type}`/`{id}

INDEXING A DOCUMENT WITH YOUR OWN ID

```
curl -X PUT
```

```
-d @part-1/first-blog-post.json
```

```
localhost:9200/blog/post/123?pretty
```

REQUEST



```
{  
  "title": "My first blog post",  
  "text": "Just trying this out...",  
  "date": "2014-01-01"  
}
```


RESPONSE



```
{
  "_index" : "blog",
  "_type" : "post",
  "_id" : "123",
  "_version" : 1,
  "_shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "created" : true
}
```

INDEXING A DOCUMENT WITH AUTOGENERATED ID

POST /{index}/{type}

** Autogenerated IDs are Base64-encoded UUIDs*

INDEXING A DOCUMENT WITH AUTOGENERATED ID

```
curl -X POST  
-d @part-1/second-blog-post.json  
localhost:9200/blog/post?pretty
```



```
{  
  "title": "Second blog post",  
  "text": "Still trying this out...",  
  "date": "2014-01-01"  
}
```

RESPONSE



```
{
  "_index" : "blog",
  "_type" : "post",
  "_id" : "AVFWIbMf7YZ6Se7RwMws",
  "_version" : 1,
  "_shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "created" : true
}
```

RETRIEVING A DOCUMENT WITH METADATA

GET `/``{index}``/``{type}``/``{id}`

RETRIEVING A DOCUMENT WITH METADATA

```
curl -X GET
```

```
localhost:9200/blog/post/123?pretty
```

RESPONSE



```
{
  "_index" : "blog",
  "_type" : "post",
  "_id" : "123",
  "_version" : 1,
  "found" : true,
  "_source": {
    "title": "My first blog entry",
    "text": "Just trying this out...",
    "date": "2014-01-01"
  }
}
```


RETRIEVING A DOCUMENT WITHOUT METADATA

GET `/``{index}``/``{type}``/``{id}``/_source`

RETRIEVING A DOCUMENT WITHOUT METADATA

```
curl -X GET
```

```
localhost:9200/blog/post/123/  
_source?pretty
```

RESPONSE



```
{  
  "title": "My first blog entry",  
  "text": "Just trying this out...",  
  "date": "2014-01-01"  
}
```

RETRIEVING PART OF A DOCUMENT

GET `/`**`{index}`**`/`**`{type}`**`/`**`{id}`**
`?_source=`**`{fields}`**

RETRIEVING PART OF A DOCUMENT

```
curl -X GET
```

```
'localhost:9200/blog/post/123?  
_source=title,date&pretty'
```

RESPONSE



```
{
  "_index" : "blog",
  "_type" : "post",
  "_id" : "123",
  "_version" : 1,
  "found" : true,
  "_source": {
    "title": "My first blog entry",
    "date": "2014-01-01"
  }
}
```

CHECKING WHETHER A DOCUMENT EXISTS

HEAD /{index}/{type}/{id}

CHECKING WHETHER A DOCUMENT EXISTS

```
curl -i -X HEAD
```

```
localhost:9200/blog/post/123
```


RESPONSE



HTTP/1.1 200 OK

Content-Length: 0

CHECKING WHETHER A DOCUMENT EXISTS

```
curl -i -X HEAD
```

```
localhost:9200/blog/post/666
```

RESPONSE



HTTP/1.1 404 Not Found

Content-Length: 0

UPDATING A WHOLE DOCUMENT

PUT `/`{index}`/`{type}`/`{id}

UPDATING A WHOLE DOCUMENT

```
curl -X PUT
```

```
-d @part-1/updated-blog-post.json
```

```
localhost:9200/blog/post/123?pretty
```

REQUEST



```
{  
  "title": "My first blog post",  
  "text": "I am starting to get the  
hang of this...",  
  "date": "2014-01-02"  
}
```

RESPONSE



```
{  
  "_index" : "blog",  
  "_type" : "post",  
  "_id" : "123",  
  "_version" : 2,  
  "_shards" : {  
    "total" : 2,  
    "successful" : 1,  
    "failed" : 0  
  },  
  "created" : false  
}
```

DELETING A DOCUMENT

DELETE /{index}/{type}/{id}

DELETING A DOCUMENT

```
curl -X DELETE
```

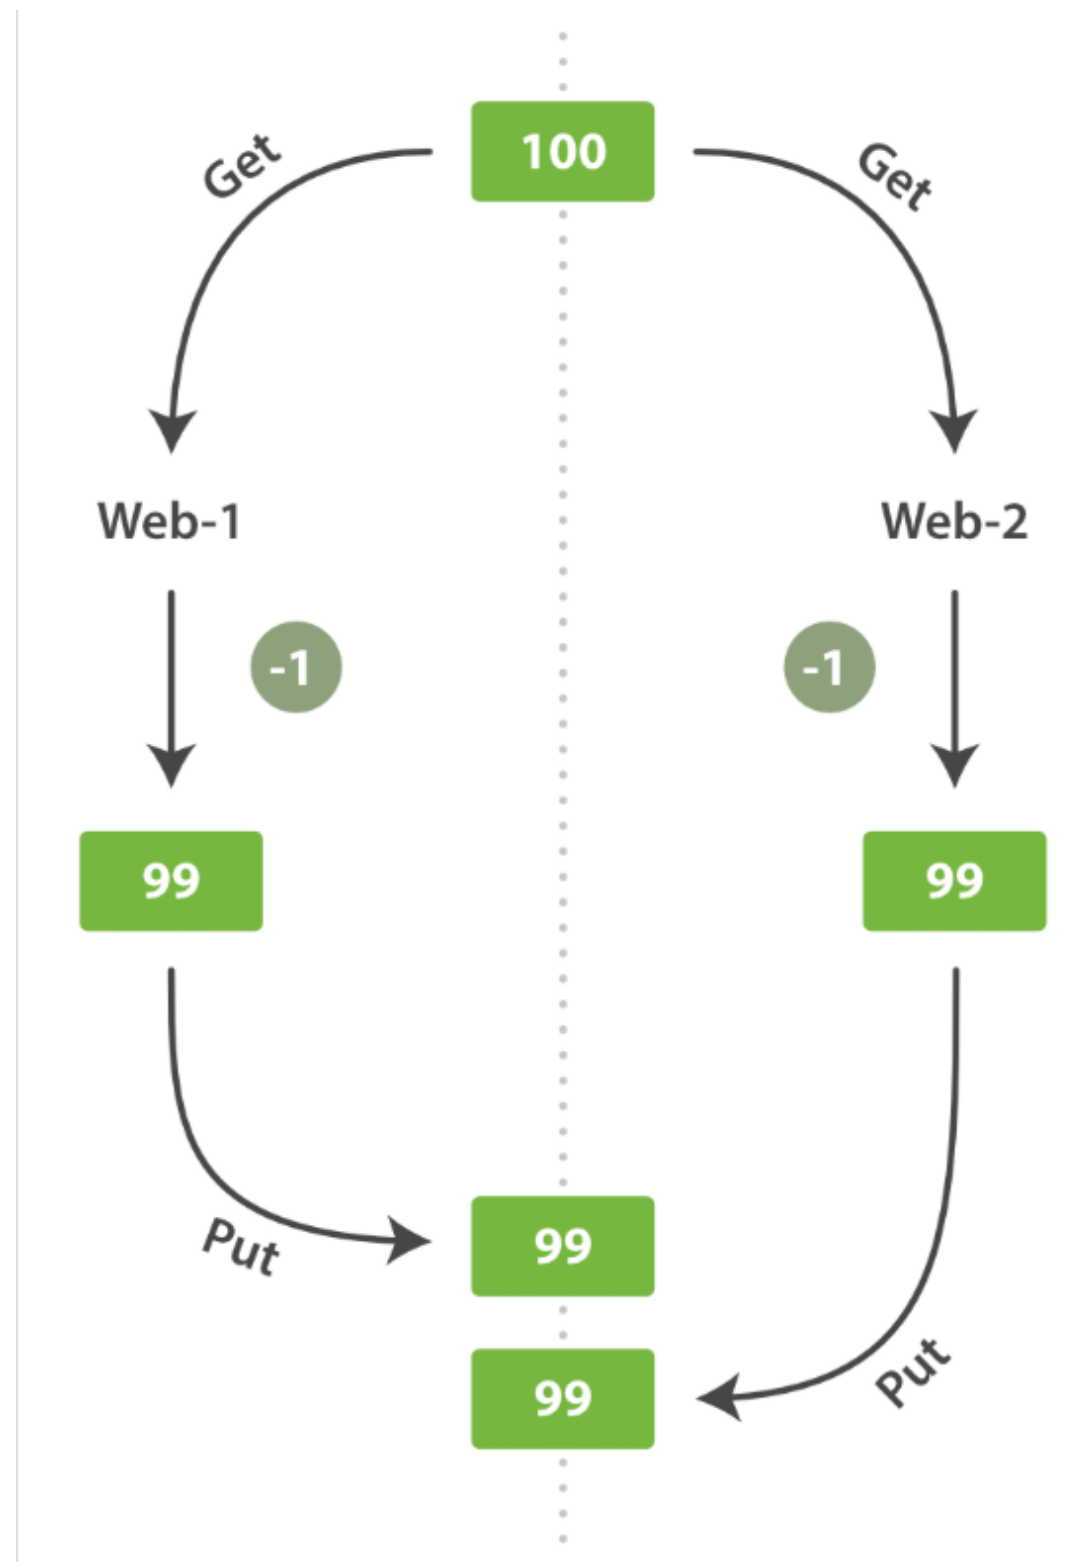
```
localhost:9200/blog/post/123?pretty
```

RESPONSE



```
{  
  "found" : true,  
  "_index" : "blog",  
  "_type" : "post",  
  "_id" : "123",  
  "_version" : 3,  
  "_shards" : {  
    "total" : 2,  
    "successful" : 1,  
    "failed" : 0  
  }  
}
```

DEALING WITH CONFLICTS



PESSIMISTIC CONCURRENCY CONTROL

- ❑ Used by relational databases
- ❑ Assumes conflicts are likely to happen (*pessimist*)
- ❑ *Blocks* access to resources

OPTIMISTIC CONCURRENCY CONTROL

- ❑ Assumes conflicts are unlikely to happen (*optimist*)
- ❑ Does *not* block operations
- ❑ If conflict happens, update *fails*

HOW ELASTICSEARCH DEALS WITH CONFLICTS

- ❑ Locking distributed resources would be very *inefficient*
- ❑ Uses *Optimistic Concurrency Control*
- ❑ Auto-increments **_version** number

HOW ELASTICSEARCH DEALS WITH CONFLICTS

- ❑ PUT /blog/post/123?version=1
- ❑ If version is outdated returns *409 Conflict*

1-3 SEARCH EXAMPLES

EMPLOYEE DIRECTORY EXAMPLE

- Index: *megacorp*
- Type: *employee*
- Ex: John Smith, Jane Smith, Douglas Fir

EMPLOYEE DIRECTORY EXAMPLE

curl -X PUT

-d @part-1/john-smith.json

localhost:9200/megacorp/employee/1

REQUEST



```
{  
  "first_name": "John",  
  "last_name": "Smith",  
  "age": 25,  
  "about": "I love to go rock climbing",  
  "interests": ["sports", "music"]  
}
```

EMPLOYEE DIRECTORY EXAMPLE

curl -X PUT

-d @part-1/jane-smith.json

localhost:9200/megacorp/employee/2

REQUEST



```
{  
  "first_name": "Jane",  
  "last_name": "Smith",  
  "age": 32,  
  "about": "I like to collect rock albums",  
  "interests": ["music"]  
}
```

EMPLOYEE DIRECTORY EXAMPLE

curl -X PUT

-d @part-1/douglas-fir.json

localhost:9200/megacorp/employee/3

REQUEST



```
{  
  "first_name": "Douglas",  
  "last_name": "Fir",  
  "age": 35,  
  "about": "I like to build cabinets",  
  "interests": ["forestry"]  
}
```

SEARCHES ALL EMPLOYEES

GET /megacorp/employee/_search

SEARCHES ALL EMPLOYEES

```
curl -X GET
```

```
localhost:9200/megacorp/employee/  
_search?pretty
```

SEARCH WITH QUERY-STRING

GET /megacorp/employee/_search
?q=last_name:Smith

SEARCH WITH QUERY-STRING

```
curl -X GET
```

```
'localhost:9200/megacorp/employee/  
_search?q=last_name:Smith&pretty'
```

RESPONSE



```
"hits" : {
  "total" : 2,
  "max_score" : 0.30685282,
  "hits" : [ {
    ...
    "_score" : 0.30685282,
    "_source" : {
      "first_name" : "Jane",
      "last_name" : "Smith", ... }
  }, {
    ...
    "_score" : 0.30685282,
    "_source" : {
      "first_name" : "John",
      "last_name" : "Smith", ... }
  } ]
}
```

SEARCH WITH QUERY DSL

```
curl -X GET
```

```
-d @part-1/last-name-query.json
```

```
localhost:9200/megacorp/employee/  
_search?pretty
```

REQUEST



```
{  
  "query": {  
    "match": {  
      "last_name": "Smith"  
    }  
  }  
}
```

RESPONSE



```
"hits" : {
  "total" : 2,
  "max_score" : 0.30685282,
  "hits" : [ {
    ...
    "_score" : 0.30685282,
    "_source" : {
      "first_name" : "Jane",
      "last_name" : "Smith", ... }
  }, {
    ...
    "_score" : 0.30685282,
    "_source" : {
      "first_name" : "John",
      "last_name" : "Smith", ... }
  } ]
}
```

SEARCH WITH QUERY DSL AND FILTER

```
curl -X GET  
-d @part-1/last-name-age-query.json  
localhost:9200/megacorp/employee/  
_search?pretty
```


REQUEST



```
"query": {  
  "filtered": {  
    "filter": {  
      "range": {  
        "age": { "gt": 30 }  
      }  
    },  
    "query": {  
      "match": { "last_name": "Smith" }  
    }  
  }  
}
```

RESPONSE



```
"hits" : {  
  "total" : 1,  
  "max_score" : 0.30685282,  
  "hits" : [ {  
    ...  
    "_score" : 0.30685282,  
    "_source": {  
      "first_name": "Jane",  
      "last_name": "Smith",  
      "age": 32, ... }  
  ]  
}
```

FULL-TEXT SEARCH

```
curl -X GET
```

```
-d @part-1/full-text-search.json
```

```
localhost:9200/megacorp/employee/  
_search?pretty
```

REQUEST



```
{  
  "query": {  
    "match": {  
      "about": "rock climbing"  
    }  
  }  
}
```

RESPONSE



```
"hits" : [{ ...
  "_score" : 0.16273327,
  "_source": {
    "first_name": "John", "last_name": "Smith",
    "about": "I love to go rock climbing", ... }
}, { ...
  "_score" : 0.016878016,
  "_source": {
    "first_name": "Jane", "last_name": "Smith",
    "about": "I like to collect rock albums", ... }
}]
```

RELEVANCE SCORES

- The **_score** field ranks searches results
- The *higher* the score, the *better*

PHRASE SEARCH

```
curl -X GET
```

```
-d @part-1/phrase-search.json
```

```
localhost:9200/megacorp/employee/  
_search?pretty
```

REQUEST



```
{  
  "query": {  
    "match_phrase": {  
      "about": "rock climbing"  
    }  
  }  
}
```


RESPONSE




```
"hits" : {  
  "total" : 1,  
  "max_score" : 0.23013961,  
  "hits" : [ {  
    ...  
    "_score" : 0.23013961,  
    "_source": {  
      "first_name": "John",  
      "last_name": "Smith",  
      "about": "I love to go rock climbing"  
      ... }  
    } ]  
  }
```

1-4 DATA RESILIENCY

CALL ME MAYBE


- *Jepsen Tests*
- Simulates *network partition* scenarios
- Run several operations against a distributed system
- Verify that the history of those operations makes sense

NETWORK PARTITION



ELASTICSEARCH STATUS

- ❑ Risk of data loss on network partition and *split-brain* scenarios



IT IS NOT SO BAD...

- ❑ Still much more resilient than *MongoDB*
- ❑ Elastic is working hard to improve it
- ❑ *Two-phase commits* are planned

IF YOU REALLY CARE ABOUT YOUR DATA

- ❑ Use a more reliable *primary* data store:
 - ❑ Cassandra
 - ❑ Postgres
- ❑ *Synchronize* it to Elasticsearch
- ❑ ...or set-up comprehensive *back-up*

*There's no such thing as a 100%
reliable distributed system*

1-5 SOLR COMPARISON

- *SolrCloud*

- Both:

 - Are *open-source* and *mature*

 - Are based on *Apache Lucene*

 - Have more or less *similar features*

SOLR API

- ❑ *HTTP GET*
- ❑ Query parameters passed in as URL parameters
- ❑ Is *not* RESTful
- ❑ Multiple formats (JSON, XML...)

SOLR API

- ❑ Version 4.4 added *Schemaless* API
- ❑ Older versions require up-front *Schema*

ELASTICSEARCH API

- *RESTful*
- *Schemaless*
- CRUD document operations
- Manage indices, read metrics, etc...

ELASTICSEARCH API

- ❑ *Query DSL*
- ❑ Better *readability*
- ❑ JSON-only

SEARCH

- Both are very good with *text search*
- Both based on *Apache Lucene*

EASYNES OF USE

- ❑ Elasticsearch is *simpler*:
 - ❑ Just a single process
 - ❑ Easier API
- ❑ SolrCloud requires *Apache ZooKeeper*

SOLRCLOUD DATA RESILIENCY

- ❑ SolrCloud uses Apache ZooKeeper to discover nodes
- ❑ Better at preventing *split-brain* conditions
- ❑ *Jepsen Tests* pass

ANALYTICS

- ❑ Elasticsearch is the choice for *analytics*:
 - ❑ Comprehensive aggregations
 - ❑ Thousands of metrics
 - ❑ SolrCloud is not even close

ThoughtWorks®

PART 2

Search and Analytics

2-1 SEARCH

Finding the needle in the haystack

TWEETS EXAMPLE

□ /<country_code>/user

□ /<country_code>/tweet

TWEETS EXAMPLE

/us/user/1

```
{  
  "email": "john@smith.com",  
  "name": "John Smith",  
  "username": "@john"  
}
```

TWEETS EXAMPLE

/gb/user/2

```
{  
  "email": "mary@jones.com",  
  "name": "Mary Jones",  
  "username": "@mary"  
}
```

TWEET EXAMPLE

/gb/tweet/3

```
{  
  "date": "2014-09-13",  
  "name": "Mary Jones",  
  "tweet": "Elasticsearch means full  
text search has never been so easy",  
  "user_id": 2  
}
```


TWEETS EXAMPLE

`./part-2/load-tweet-data.sh`

THE EMPTY SEARCH

- Returns all documents on all indices

GET /_search

THE EMPTY SEARCH

```
curl -X GET
```

```
localhost:9200/_search?pretty
```

THE EMPTY SEARCH

```
"hits" : {  
  "total" : 14,  
  "hits" : [  
    {  
      "_index": "us",  
      "_type": "tweet",  
      "_id": "7",  
      "_score": 1,  
      "_source": {  
        "date": "2014-09-17",  
        "name": "John Smith",  
        "tweet": "The Query DSL is really powerful and flexible",  
        "user_id": 2  
      }  
    },  
    ... 9 RESULTS REMOVED ...  
  ]  
}
```

MULTI-INDEX, MULTITYPE SEARCH

□ /_search

□ /gb/_search

□ /gb,us/_search

□ /gb/user/_search

□ /_all/user,tweet/_search

PAGINATION

- ❑ Returns *10 results* per request (default)
- ❑ Control parameters:
 - ❑ *size*: number of results to return
 - ❑ *from*: number of results to skip

PAGINATION

- GET `/_search?size=5`
- GET `/_search?size=5&from=5`
- GET `/_search?size=5&from=10`

TYPES OF SEARCH

- ❑ *Structured query* on concrete fields (similar to SQL)
- ❑ *Full-text query* (sorts results by relevance)
- ❑ Combination of the two

SEARCH BY EXACT VALUES

- Examples:
 - date
 - user ID
 - username
- “Does this document *match* the query?”

SEARCH BY EXACT VALUES

□ SQL queries:

```
SELECT * FROM user
WHERE name = "John Smith"
      AND user_id = 2
      AND date > "2014-09-15"
```

FULL-TEXT SEARCH

- Examples:
 - the text of a tweet
 - body of an email
- “*How well* does this document match the query?”

FULL-TEXT SEARCH

- ❑ *UK* should also match *United Kingdom*
- ❑ *jump* should also match *jumped, jumps, jumping* and *leap*

FULL-TEXT SEARCH

- ❑ *fox news hunting* should return stories about *hunting on Fox News*
- ❑ *fox hunting news* should return news stories about *fox hunting*

HOW ELASTICSEARCH PERFORMS TEXT SEARCH

- *Analyzes* the text
 - *Tokenizes* into terms
 - *Normalizes* the terms
- Builds an *inverted index*

LIST OF INDEXED DOCUMENTS

ID	Text
1	Baseball is played during summer months.
2	Summer is the time for picnics here.
3	Months later we found out why.
4	Why is summer so hot here.

INVERTED INDEX

Term	Frequency	Document IDs
<i>baseball</i>	1	1
<i>during</i>	1	1
<i>found</i>	1	3
<i>here</i>	2	2, 4
<i>hot</i>	1	4
<i>is</i>	3	1, 2, 4
<i>months</i>	2	1, 3
<i>summer</i>	3	1, 2, 4
<i>the</i>	1	2
<i>why</i>	2	3, 4

QUERY DSL

GET /_search

{

"query": *YOUR_QUERY_HERE*

}

QUERY BY FIELD

```
{  
  "match": {  
    "tweet": "elasticsearch"  
  }  
}
```

QUERY BY FIELD

```
curl -X GET -d
```

```
@part-2/elasticsearch-tweets-query.json
```

```
localhost:9200/_all/tweet/_search
```

QUERY WITH MULTIPLE CLAUSES

```
{
  "bool": {
    "must": {
      "match": { "tweet": "elasticsearch" }
    },
    "must_not": {
      "match": { "name": "mary" }
    },
    "should": {
      "match": { "tweet": "full text" }
    }
  }
}
```

QUERY WITH MULTIPLE CLAUSES

```
curl -X GET -d
```

```
@part-2/combining-tweet-queries.json
```

```
localhost:9200/_all/tweet/_search
```

QUERY WITH MULTIPLE CLAUSES

```
"_score": 0.07082729, "_source": { ...  
  "name": "John Smith",  
  "tweet": "The Elasticsearch API is really easy to use"  
}, ...
```

```
"_score": 0.049890988, "_source": { ...  
  "name": "John Smith",  
  "tweet": "Elasticsearch surely is one of the hottest  
new NoSQL products"  
}, ...
```

```
"_score": 0.03991279, "_source": { ...  
  "name": "John Smith",  
  "tweet": "Elasticsearch and I have left the honeymoon  
stage, and I still love her." }
```

MOST IMPORTANT QUERIES

- *match*
- *match_all*
- *multi_match*
- *bool*

QUERIES VS. FILTERS

- *Queries:*

- full-text

- “*how well* does the document match?”

- *Filters:*

- exact values

- *yes-no* questions

QUERIES VS. FILTERS

- The goal of *filters* is to reduce the number of documents that have to be examined by a *query*

PERFORMANCE COMPARISON

- ❑ Filters are easy to *cache* and can be reused efficiently
- ❑ Queries are heavier and *non-cacheable*

WHEN TO USE WHICH

- ❑ Use queries only for *full-text search*
- ❑ Use filters for anything else

FILTER BY EXACT FIELD VALUES

```
"filtered": {  
  "filter": {  
    "term": {  
      "user_id": 1  
    }  
  }  
}
```

FILTER BY EXACT FIELD VALUES

```
curl -X GET -d  
@part-2/user-id-filter.json  
localhost:9200/_search
```

FILTER BY EXACT FIELD VALUES

```
"filtered": {  
  "filter": {  
    "range": {  
      "date": {  
        "gte": "2014-09-20"  
      }  
    }  
  }  
}
```

FILTER BY EXACT FIELD VALUES

```
curl -X GET -d  
@part-2/date-filter.json  
localhost:9200/_search
```

MOST IMPORTANT FILTERS

- ☐ *term*
- ☐ *terms*
- ☐ *range*
- ☐ *exists* and *missing*
- ☐ *bool*

COMBINING QUERIES WITH FILTERS

```
"filtered": {  
  "query": {  
    "match": {  
      "tweet": "elasticsearch"  
    }  
  },  
  "filter": {  
    "term": { "user_id": 1 }  
  }  
}
```

COMBINING QUERIES WITH FILTERS

```
curl -X GET -d
```

```
@part-2/filtered-tweet-query.json
```

```
localhost:9200/_search
```

SORTING

- *Relevance score*
- The *higher* the score, the *better*
- By default, results are returned in *descending* order of relevance
- You can sort by *any field*

RELEVANCE SCORE

- *Similarity algorithm*
 - Term Frequency / Inverse Document Frequency (*TF/IDF*)

RELEVANCE SCORE

☐ *Term frequency*

- ☐ How often does the term appear in the field?
- ☐ The more often, the *more* relevant

RELEVANCE SCORE

- *Inverse document frequency*
 - How often does each term appear in the index?
 - The more often, the *less* relevant

RELEVANCE SCORE

□ *Field-length norm*

□ How long is the field?

□ The longer it is, the *less likely* it is that words in the field will be relevant

2-2 ANALYTICS

How many needles are in the haystack?

SEARCH

□ Just looks for the *needle* in the haystack

BUSINESS QUESTIONS

- ❑ How *many* needles are in the haystack?
- ❑ What is the needle *average length*?
- ❑ What is the *median length* of the needles, *by manufacturer*?
- ❑ How many needles were added to the haystack *each month*?

BUSINESS QUESTIONS

- ❑ What are your *most popular* needle manufactures?
- ❑ Are there any *anomalous* clumps of needles?

AGGREGATIONS

- ❑ Answer *Analytics* questions
- ❑ Can be combined with *Search*
- ❑ Near *real-time* in Elasticsearch
 - ❑ SQL queries can take *days*

AGGREGATIONS

Buckets + Metrics

BUCKETS

- ❑ Collection of documents that meet a certain *criteria*
- ❑ Can be *nested* inside other buckets

BUCKETS

- *Employee* \Rightarrow *male* or *female* bucket
- *San Francisco* \Rightarrow *California* bucket
- *2014-10-28* \Rightarrow *October* bucket

METRICS

- *Calculations* on top of buckets
- Answer the questions
- Ex: *min, max, mean, sum...*

EXAMPLE

- ❑ Partition by *country* (*bucket*)
- ❑ ...then partition by *gender* (*bucket*)
- ❑ ...then partition by *age ranges* (*bucket*)
- ❑ ...calculate the *average salary* for each age range (*metric*)

CAR TRANSACTIONS EXAMPLE

□ /cars/transactions

CAR TRANSACTIONS EXAMPLE

/cars/transactions/
AVFr1xbVmdUYWpF46Ps4

```
{  
  "price" : 10000,  
  "color" : "red",  
  "make" : "honda",  
  "sold" : "2014-10-28"  
}
```

CAR TRANSACTIONS EXAMPLE

`./part-2/load-car-data.sh`

BEST SELLING CAR COLOR

```
{  
  "aggs": {  
    "colors": {  
      "terms": {  
        "fields": "color"  
      }  
    }  
  }  
}
```

BEST SELLING CAR COLOR

```
curl -X GET -d  
@part-2/best-selling-car-color.json  
'localhost:9200/cars/transactions/  
_search?search_type=count&pretty'
```

BEST SELLING CAR COLOR

```
"colors" : {  
  "buckets" : [{  
    "key" : "red",  
    "doc_count" : 16  
  }, {  
    "key" : "blue",  
    "doc_count" : 8  
  }, {  
    "key" : "green",  
    "doc_count" : 8  
  }]  
}
```

AVERAGE CAR COLOR PRICE

```
{
  "aggs": {
    "colors": {
      "terms": { "field": "color" },
      "aggs": {
        "avg_price": {
          "avg": { "field": "price" }
        }
      }
    }
  }
}
```


AVERAGE CAR COLOR PRICE

```
curl -X GET -d  
@part-2/average-car-color-price.json  
'localhost:9200/cars/transactions/  
_search?search_type=count&pretty'
```

AVERAGE CAR COLOR PRICE

```
"colors" : {  
  "buckets": [{  
    "key": "red", "doc_count": 16,  
    "avg_price": { "value": 32500.0 }  
  }, {  
    "key": "blue", "doc_count": 8,  
    "avg_price": { "value": 20000.0 }  
  }, {  
    "key": "green", "doc_count": 8,  
    "avg_price": { "value": 21000.0 }  
  }]  
}
```

BUILDING BAR CHARTS

- ❑ Very easy to convert aggregations to *charts* and *graphs*
- ❑ Ex: *histograms* and *time-series*

CAR SALES REVENUE HISTOGRAM

```
{
  "aggs": {
    "price": {
      "histogram": {
        "field": "price",
        "interval": 20000
      },
      "aggs": {
        "revenue": {"sum": {"field": "price"}}
      }
    }
  }
}
```


CAR SALES REVENUE HISTOGRAM

```
curl -X GET -d  
@part-2/car-revenue-histogram.json  
'localhost:9200/cars/transactions/  
_search?search_type=count&pretty'
```

CAR SALES REVENUE HISTOGRAM

```
"price" : {  
  "buckets": [  
    { "key": 0, "doc_count": 12,  
      "revenue": {"value": 148000.0} },  
    { "key": 20000, "doc_count": 16,  
      "revenue": {"value": 380000.0} },  
    { "key": 40000, "doc_count": 0,  
      "revenue": {"value": 0.0} },  
    { "key": 60000, "doc_count": 0,  
      "revenue": {"value": 0.0} },  
    { "key": 80000, "doc_count": 4,  
      "revenue": {"value": 320000.0} }  
  ]  
}
```

CAR SALES REVENUE HISTOGRAM



TIME-SERIES DATA

- ❑ Data with a *timestamp*:
 - ❑ How many cars sold *each month* this year?
 - ❑ What was the price of this stock for the *last 12 hours*?
 - ❑ What was the average latency of our website *every hour* in the last week?

HOW MANY CARS SOLD PER MONTH?

```
{
  "aggs": {
    "sales": {
      "date_histogram": {
        "field": "sold",
        "interval": "month",
        "format": "yyyy-MM-dd"
      }
    }
  }
}
```


HOW MANY CARS SOLD PER MONTH?

```
curl -X GET -d  
@part-2/car-sales-per-month.json  
'localhost:9200/cars/transactions/  
_search?search_type=count&pretty'
```

HOW MANY CARS SOLD PER MONTH?

```
"sales" : {  
  "buckets" : [  
    {"key_as_string": "2014-01-01", "doc_count": 4},  
    {"key_as_string": "2014-02-01", "doc_count": 4},  
    {"key_as_string": "2014-03-01", "doc_count": 0},  
    {"key_as_string": "2014-04-01", "doc_count": 0},  
    {"key_as_string": "2014-05-01", "doc_count": 4},  
    {"key_as_string": "2014-06-01", "doc_count": 0},  
    {"key_as_string": "2014-07-01", "doc_count": 4},  
    {"key_as_string": "2014-08-01", "doc_count": 4},  
    {"key_as_string": "2014-09-01", "doc_count": 0},  
    {"key_as_string": "2014-10-01", "doc_count": 4},  
    {"key_as_string": "2014-11-01", "doc_count": 8}  
  ]  
}
```

HOW MANY CARS SOLD PER MONTH?



ThoughtWorks®

PART 3

Dealing with human language

3-1 INVERTED INDEX

INVERTED INDEX

- Data structure
- Efficient *full-text* search

EXAMPLE

Document 1

The quick brown fox jumped
over the lazy dog

Document 2

Quick brown foxes leap over
lazy dogs in summer

TOKENIZATION

Document 1

```
["The", "quick", "brown",  
"fox", "jumped", "over",  
"the", "lazy", "dog"]
```

Document 2

```
["Quick", "brown", "foxes",  
"leap", "over", "lazy",  
"dogs", "in", "summer"]
```

Term	Document 1	Document 2
<i>Quick</i>		×
<i>The</i>	×	
<i>brown</i>	×	×
<i>dog</i>	×	
<i>dogs</i>		×
<i>fox</i>	×	
<i>foxes</i>		×
<i>in</i>		×
<i>jumped</i>	×	
<i>lazy</i>	×	×
<i>leap</i>		×
<i>over</i>	×	×
<i>quick</i>	×	
<i>summer</i>		×
<i>the</i>	×	

EXAMPLE

□ Searching for “*quick brown*”

Term	Document 1	Document 2
<i>brown</i>	×	×
<i>quick</i>	×	
Total	2	1

□ *Naive similarity* algorithm:

□ Document 1 is a *better match*

A FEW PROBLEMS

- *Quick* and *quick* are the same word
- *fox* and *foxes* are pretty similar
- *jumped* and *leap* are synonyms

NORMALIZATION

- ❑ *Quick lowercased* to *quick*
- ❑ *foxes stemmed* to *fox*
- ❑ *jumped* and *leap replaced* by *jump*

BETTER INVERTED INDEX

Term	Document 1	Document 2
<i>brown</i>	×	×
<i>dog</i>	×	×
<i>fox</i>	×	×
<i>in</i>		×
<i>jump</i>	×	×
<i>lazy</i>	×	×
<i>over</i>	×	×
<i>quick</i>	×	×
<i>summer</i>		×
<i>the</i>	×	×

SEARCH INPUT

- ❑ You can only find terms that exist in the *inverted index*
- ❑ The *query string* is also *normalized*

3-2 ANALYZERS

ANALYSIS

- ❑ *Tokenizes* a block of text into *terms*
- ❑ *Normalizes* terms to standard form
- ❑ Improves searchability

ANALYZERS

- Pipeline:
 - Character filters
 - Tokenizer
 - Token filters

BUILT-IN ANALYZERS

- ❑ *Standard* analyzer
- ❑ *Language-specific* analyzers
 - ❑ 30+ languages supported

TESTING THE STANDARD ANALYZER

```
GET /_analyze?  
analyzer=standard
```

```
The quick brown fox jumped  
over the lazy dog.
```

TESTING THE STANDARD ANALYZER

```
curl -X GET -d  
@part-3/quick-brown-fox.txt  
'localhost:9200/_analyze?  
analyzer=standard&pretty'
```

TESTING THE STANDARD ANALYZER

```
"tokens" : [  
  {"token": "the", ...},  
  {"token": "quick", ...},  
  {"token": "brown", ...},  
  {"token": "fox", ...},  
  {"token": "jumps", ...},  
  {"token": "over", ...},  
  {"token": "the", ...},  
  {"token": "lazy", ...},  
  {"token": "dog", ...}  
]
```

TESTING THE ENGLISH ANALYZER

GET /_analyze?analyzer=english

The quick brown fox jumped
over the lazy dog.

TESTING THE ENGLISH ANALYZER

```
curl -X GET -d  
@part-3/quick-brown-fox.txt  
'localhost:9200/_analyze?  
analyzer=english&pretty'
```


TESTING THE ENGLISH ANALYZER

```
"tokens" : [  
  {"token": "quick", ...},  
  {"token": "brown", ...},  
  {"token": "fox", ...},  
  {"token": "jump", ...},  
  {"token": "over", ...},  
  {"token": "lazi", ...},  
  {"token": "dog", ...}  
]
```

TESTING THE BRAZILIAN ANALYZER

```
GET /_analyze?  
analyzer=brazilian
```

A rápida raposa marrom pulou
sobre o cachorro preguiçoso.

TESTING THE BRAZILIAN ANALYZER

```
curl -X GET -d  
@part-3/raposa-rapida.txt  
'localhost:9200/_analyze?  
analyzer=brazilian&pretty'
```

TESTING THE BRAZILIAN ANALYZER

```
"tokens" : [  
  {"token": "rap", ...},  
  {"token": "rapos", ...},  
  {"token": "marrom", ...},  
  {"token": "pul", ...},  
  {"token": "cachorr", ...},  
  {"token": "preguic", ...}  
]
```

STEMMERS

- *Algorithmic* stemmers:

- Faster

- Less precise

- *Dictionary* stemmers:

- Slower

- More precise

3-3 MAPPING

MAPPING

- ❑ Every document has a *type*
- ❑ Every type has its own *mapping*
- ❑ A mapping defines:
 - ❑ The *fields*
 - ❑ The *datatype* for each field

MAPPING

- ❑ Elasticsearch *guesses* the mapping when a new field is added
- ❑ Should customize the mapping for improved *search* and *performance*
- ❑ Must customize the mapping when *type* is created

MAPPING

- ❑ A field's mapping *cannot be changed*
- ❑ You can still add *new fields*
- ❑ Only option is to *reindex* all documents
- ❑ Reindexing with zero-downtime:
 - ❑ *index aliases*

CORE FIELD TYPES

- ❑ String
- ❑ Integer
- ❑ Floating-point
- ❑ Boolean
- ❑ Date
- ❑ *Inner Objects*

VIEWING THE MAPPING

GET `/ {index} /_mapping/ {type}`

VIEWING THE MAPPING

```
curl -X GET
```

```
'localhost:9200/gb/_mapping/  
tweet?pretty'
```

VIEWING THE MAPPING

```
"date": {  
  "type": "date",  
  "format":  
    "strict_date_optional_time..."  
},  
"name": {  
  "type": "string"  
},  
"tweet": {  
  "type": "string"  
},  
"user_id": {  
  "type": "long"  
}
```

CUSTOMIZING FIELD MAPPINGS

- ❑ Distinguish between:
 - ❑ *Full-text* string fields
 - ❑ *Exact value* string fields
- ❑ Use *language-specific* analyzers

STRING MAPPING ATTRIBUTES

- index:

- analyzed (*full-text search, default*)

- not_analyzed (*exact value*)

- analyzer:

- standard (*default*)

- english

- ...

ADDING NEW SEARCHABLE FIELD

PUT /gb,us/_mapping/tweet

```
{
  "properties": {
    "description": {
      "type": "string",
      "index": "analyzed",
      "analyzer": "english"
    }
  }
}
```


ADDING NEW SEARCHABLE FIELD

```
curl -X PUT -d  
@part-3/add-new-mapping.json  
'localhost:9200/gb,us/  
_mapping/tweet?pretty'
```

ADDING NEW SEARCHABLE FIELD

```
curl -X GET  
'localhost:9200/us,gb/  
_mapping/tweet?pretty'
```

ADDING NEW SEARCHABLE FIELD

```
...  
"description": {  
  "type": "string",  
  "analyzer": "english"  
}...
```

3-4 PROXIMITY MATCHING

THE PROBLEM

- *Sue ate the alligator*
- *The alligator ate Sue*
- *Sue never goes anywhere without her alligator-skin purse*

THE PROBLEM

- ❑ Search for “*sue alligator*” would match all three
- ❑ *Sue* and *alligator* may be separated by *paragraphs* of other text

HEURISTIC

- Words that appear *near each other* are probably related
- Give documents in which the words are close together a *higher relevance score*

TERM POSITIONS

GET /_analyze?
analyzer=standard

Quick brown fox.

TERM POSITIONS

```
"tokens": [  
  { "token": "quick", ...  
    "position": 1 },  
  { "token": "brown", ...  
    "position": 2 },  
  { "token": "fox", ...  
    "position": 3 }  
]
```


EXACT PHRASE MATCHING

GET /{index}/{type}/_search

```
{  
  "query": {  
    "match_phrase": {  
      "title": "quick brown fox"  
    }  
  }  
}
```

EXACT PHRASE MATCHING

- ❑ *quick*, *brown* and *fox* must all appear
- ❑ The position of *brown* must be 1 greater than the position of *quick*
- ❑ The position of *fox* must be 2 greater than the position of *quick*



FLEXIBLE PHRASE MATCHING

- ❑ Exact phrase matching is *too strict*
- ❑ “*quick fox*” should also match
- ❑ *Slop* matching



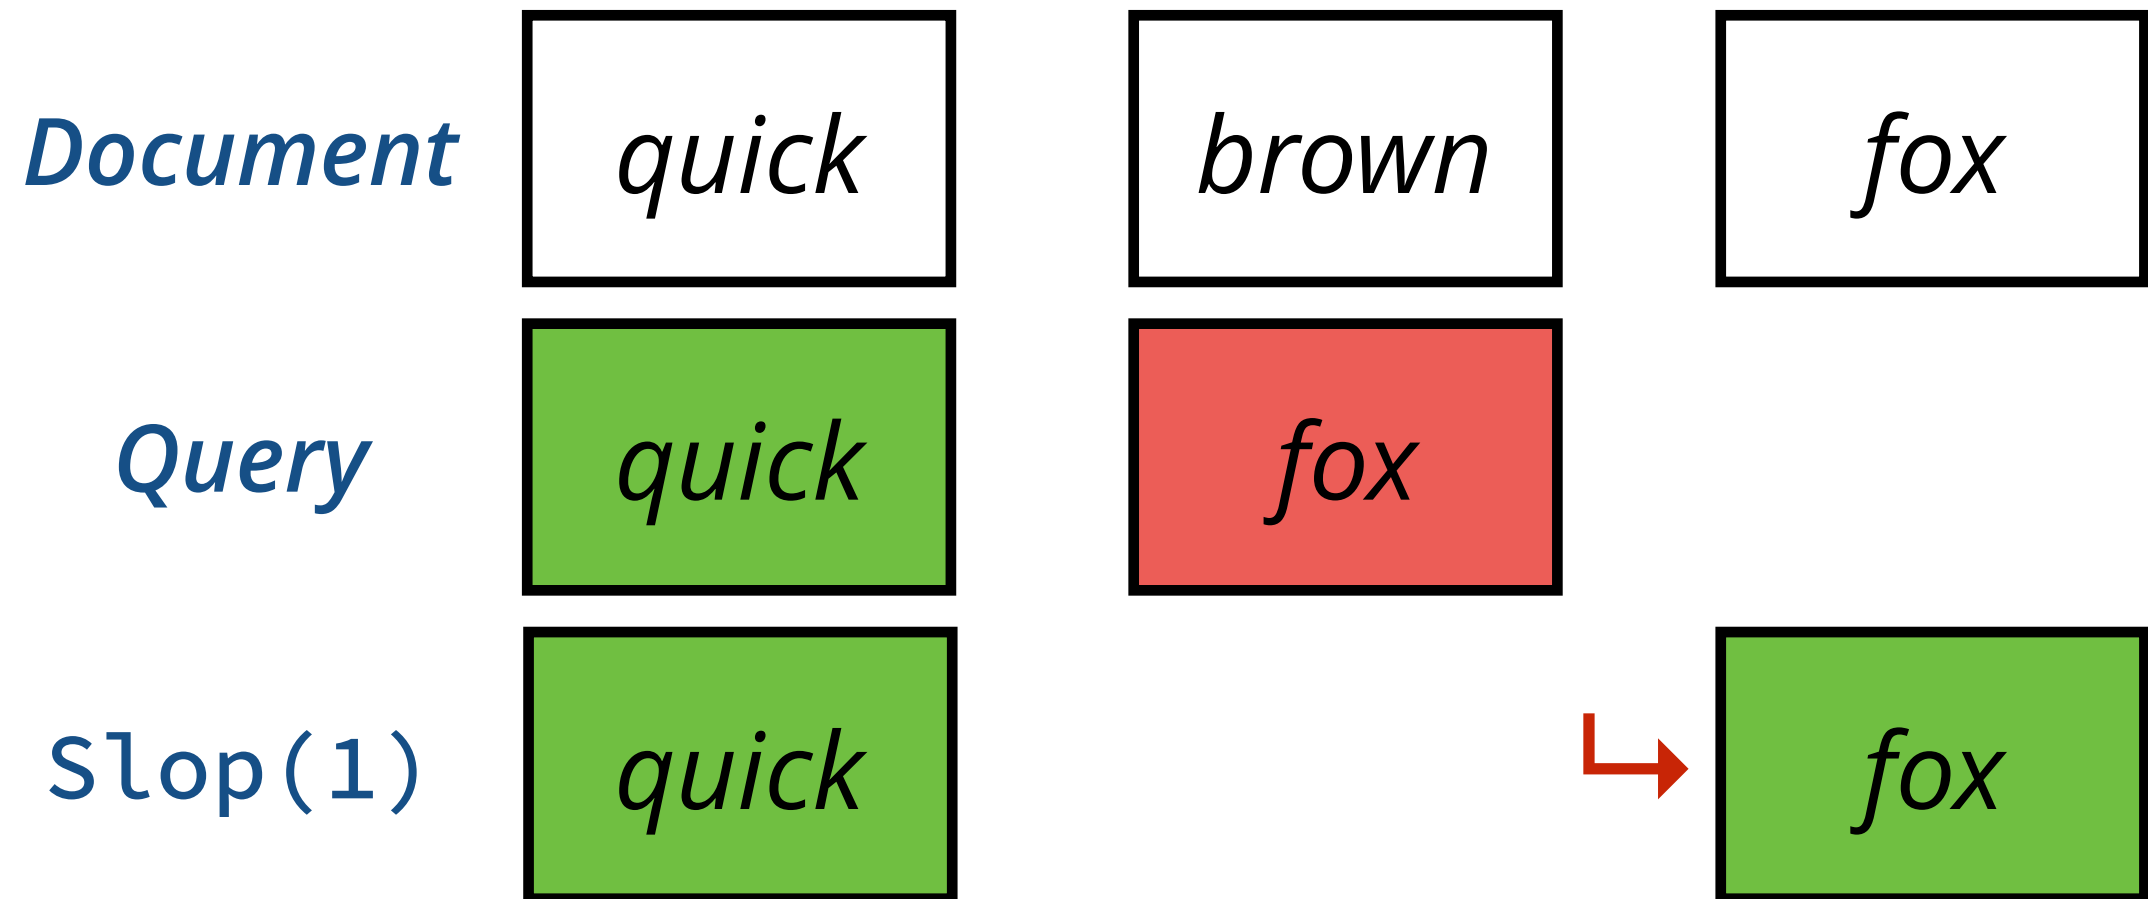
FLEXIBLE PHRASE MATCHING

```
"query": {  
  "match_phrase": {  
    "title": {  
      "query": "quick fox",  
      "slop": 1  
    }  
  }  
}
```

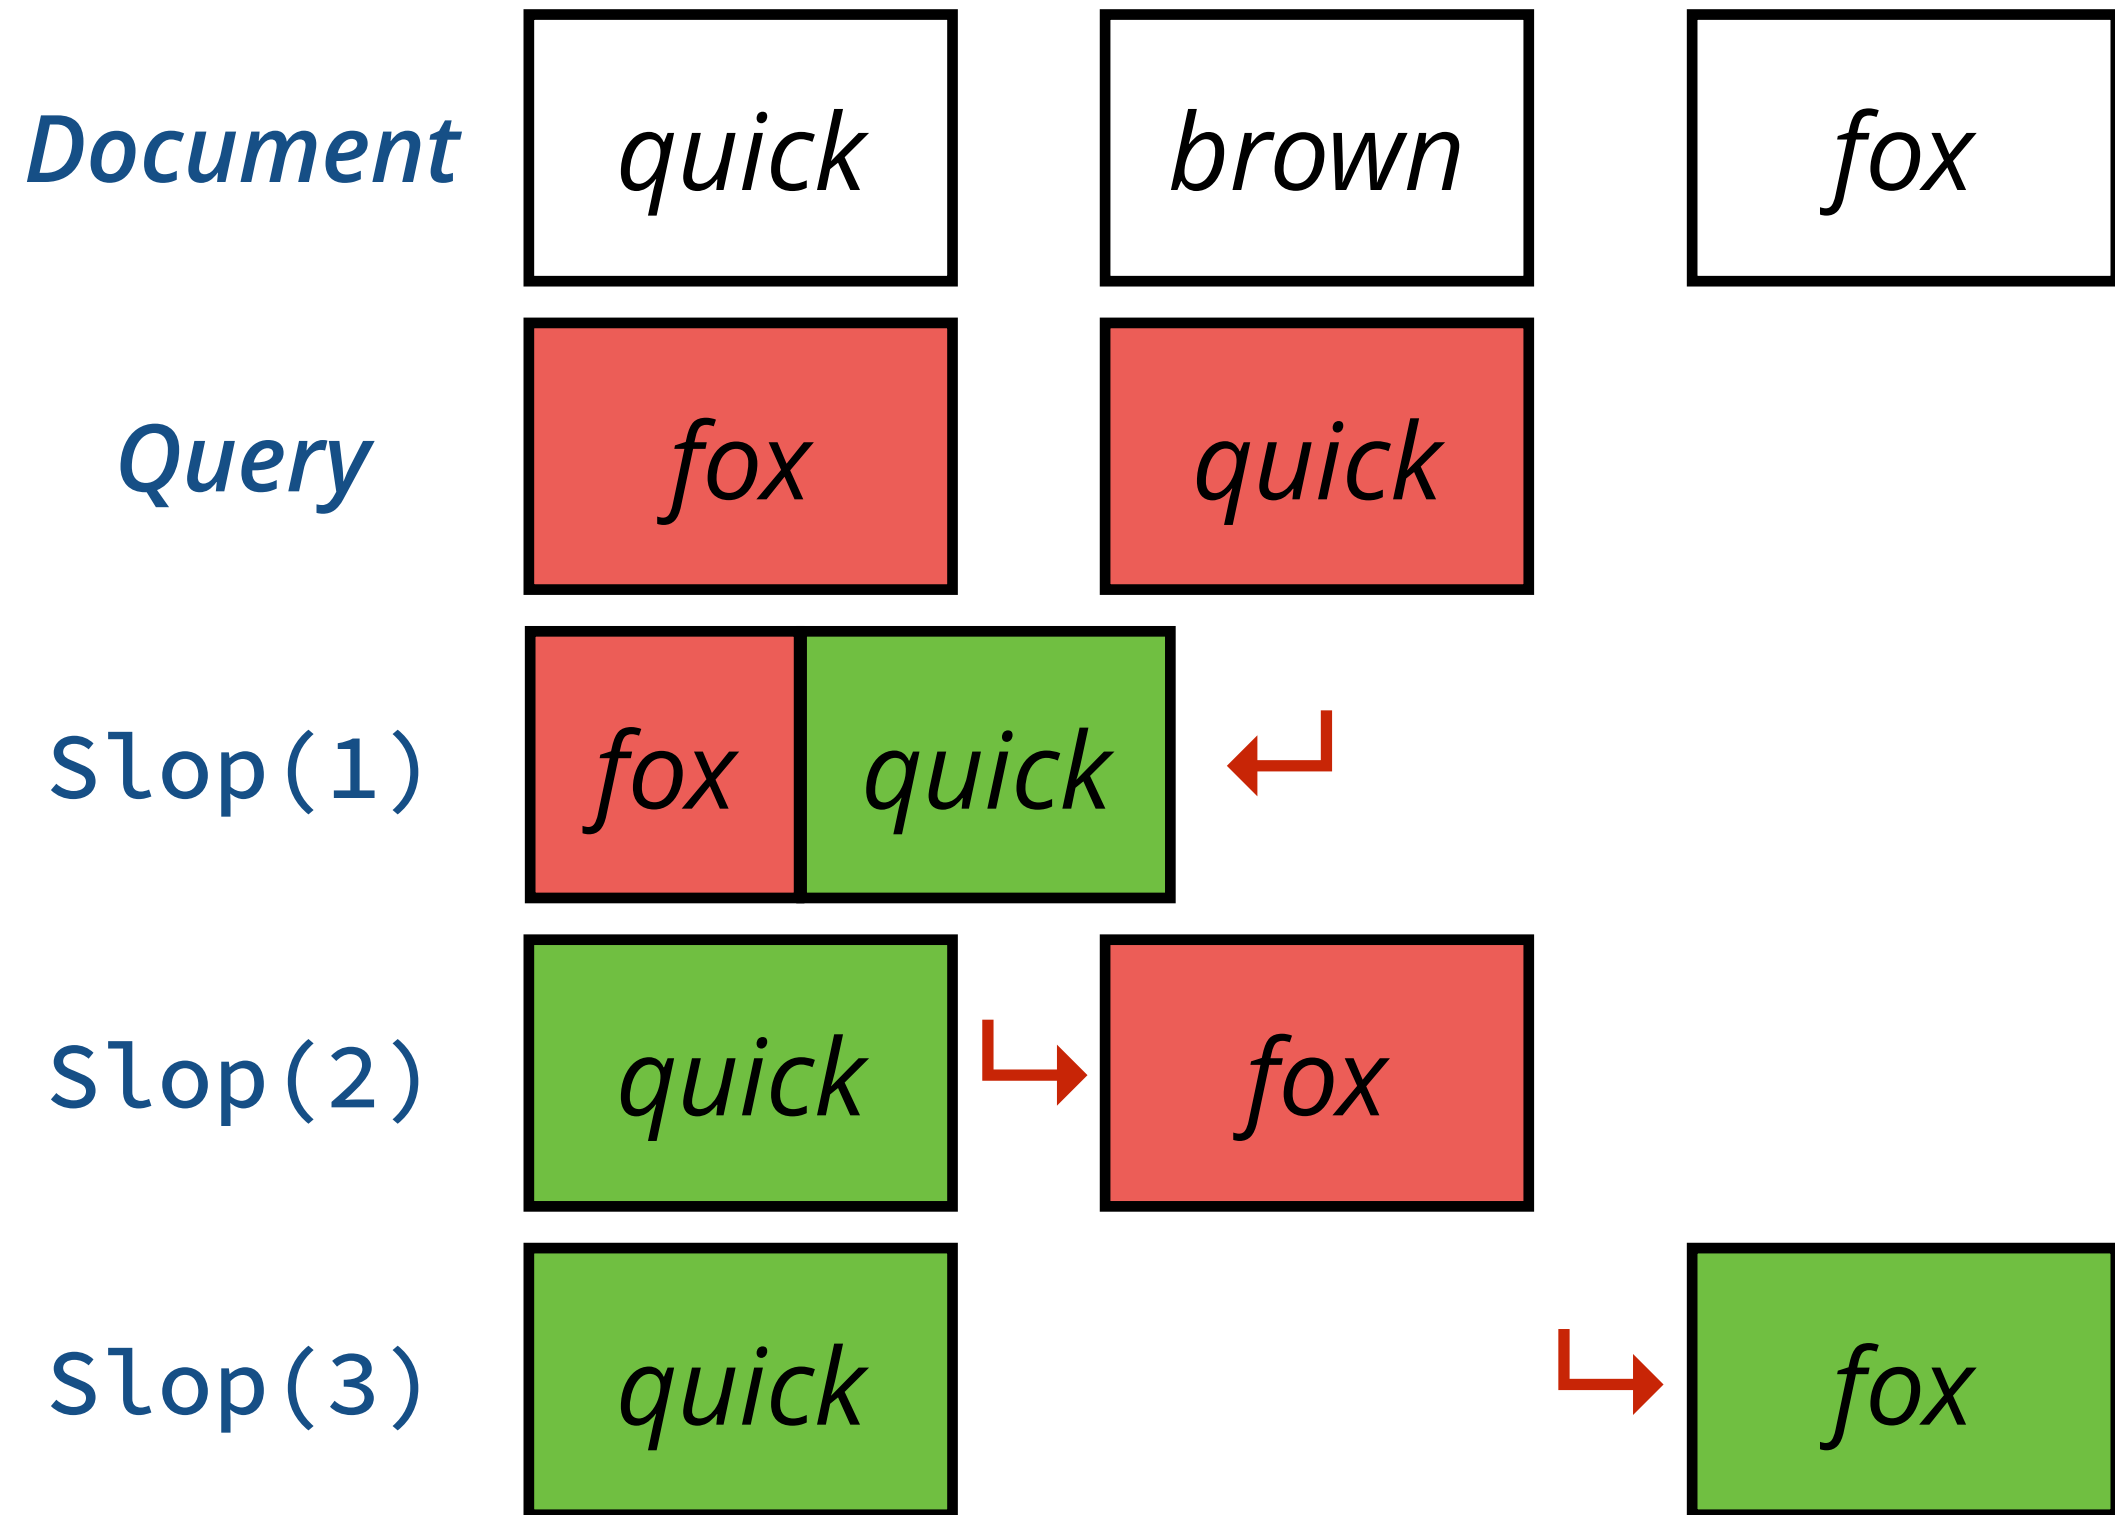
SLOP MATCHING

- How *many times* you are allowed to *move a term* in order to make the query and document match?
 - $\text{Slop}(n)$

SLOP MATCHING



SLOP MATCHING



3-5 FUZZY MATCHING

FUZZY MATCHING

- *quick brown fox* → *fast brown foxes*
- *Johnny Walker* → *Johnnie Walker*
- *Shcwarzenneger* → *Schwarzenegger*

DAMERAU-LEVENSHTEIN EDIT DISTANCE

- One-character edits:
 - *Substitution*
 - *Insertion*
 - *Deletion*
 - *Transposition* of two adjacent characters

DAMERAU-LEVENSHTEIN EDIT DISTANCE

□ One-character *substitution*:

□ fox → box

DAMERAU-LEVENSHTEIN EDIT DISTANCE

□ *Insertion* of a new character:

□ sic → sick

DAMERAU-LEVENSHTEIN EDIT DISTANCE

□ *Deletion* of a character:

□ b|ack → back

DAMERAU-LEVENSHTEIN EDIT DISTANCE

□ *Transposition* of two adjacent characters:

□ star → tsar

DAMERAU-LEVENSHTEIN EDIT DISTANCE

□ Converting *bieber* into *beaver*

1. *Substitute*: bie**b**er → bie**v**er

2. *Substitute*: bi**e**ver → ba**e**ver

3. *Transpose*: ba**e**ver → be**a**ver

□ *Edit distance* of 3

FUZZINESS

- ❑ *80% of human misspellings* have an Edit Distance of 1
- ❑ Elasticsearch supports a *maximum* Edit Distance of 2
- ❑ `fuzziness` operator

FUZZINESS EXAMPLE

`./part-3/load-surprise-data.sh`

QUERY WITHOUT FUZZINESS

GET /example/surprise/_search

```
{  
  "query": {  
    "match": {  
      "text": {  
        "query": "surprise"  
      }  
    }  
  }  
}
```

QUERY WITHOUT FUZZINESS

```
curl -X GET -d  
@part-3/surprise-query.json  
'localhost:9200/example/  
surprise/_search?pretty'
```

QUERY WITHOUT FUZZINESS

```
"hits": {  
  "total": 0,  
  "max_score": null,  
  "hits": []  
}
```

QUERY WITH FUZZINESS

GET /example/surprise/_search

```
{
  "query": {
    "match": {
      "text": {
        "query": "surprise",
        "fuzziness": "1"
      }
    }
  }
}
```

QUERY WITH FUZZINESS

```
curl -X GET -d  
@part-3/surprise-fuzzy-  
query.json  
'localhost:9200/example/  
surprise/_search?pretty'
```

QUERY WITH FUZZINESS

```
"hits": [ {  
  "_index": "example",  
  "_type": "surprise",  
  "_id": "1",  
  "_score": 0.19178301,  
  "_source": { "text": "Surprise me!" }  
}]
```


AUTO-FUZZINESS

- ❑ 0 for strings of *one* or *two* characters
- ❑ 1 for strings of *three, four* or *five* characters
- ❑ 2 for strings of *more than five* characters

ThoughtWorks®

PART 4

Data modeling

4-1 INSIDE A CLUSTER

NODES AND CLUSTERS

- ❑ A *node* is a machine running Elasticsearch
- ❑ A cluster is a set of *nodes* in the same network and with the same *cluster name*

SHARDS

- ❑ A node stores data inside its *shards*
- ❑ Shards are the *smallest unit* of *scale* and *replication*
- ❑ Each shard is a completely independent *Lucene index*

AN EMPTY CLUSTER

CLUSTER

NODE 1 - ★ MASTER

CLUSTER HEALTH

GET `/_cluster/health`

CLUSTER HEALTH

```
"cluster_name": "elasticsearch",  
"status": "green",  
"number_of_nodes": 1,  
"number_of_data_nodes": 1,  
"active_primary_shards": 0,  
"active_shards": 0,  
"relocating_shards": 0,  
"initializing_shards": 0,  
"unassigned_shards": 0
```


ADD AN INDEX

PUT /blogs

```
"settings": {  
  "number_of_shards": 3,  
  "number_of_replicas": 1  
}
```

ADD AN INDEX

CLUSTER

NODE 1 - ★ MASTER

P0

P1

P2

CLUSTER HEALTH

GET `/_cluster/health`

CLUSTER HEALTH

```
"cluster_name": "elasticsearch",  
"status": "yellow",  
"number_of_nodes": 1,  
"number_of_data_nodes": 1,  
"active_primary_shards": 3,  
"active_shards": 3,  
"relocating_shards": 0,  
"initializing_shards": 0,  
"unassigned_shards": 3
```

ADD A BACKUP NODE

CLUSTER

NODE 1 - ★ MASTER

P0

P1

P2

NODE 2

R0

R1

R2

CLUSTER HEALTH

GET `/_cluster/health`

CLUSTER HEALTH

```
"cluster_name": "elasticsearch",  
"status": "green",  
"number_of_nodes": 2,  
"number_of_data_nodes": 2,  
"active_primary_shards": 3,  
"active_shards": 6,  
"relocating_shards": 0,  
"initializing_shards": 0,  
"unassigned_shards": 0
```

THREE NODES

CLUSTER

NODE 1 - ★ MASTER



P1

P2

NODE 2

R0

R1



NODE 3

P0



R2

INCREASING THE NUMBER OF REPLICAS

PUT /blogs

```
"settings": {  
  "number_of_shards": 3,  
  "number_of_replicas": 2  
}
```

INCREASING THE NUMBER OF REPLICAS

CLUSTER

NODE 1 - ★ MASTER

R0

P1

P2

NODE 2

R0

R1

R2

NODE 3

P0

R1

R2

NODE 1 FAILS

CLUSTER

NODE 2 - ★ MASTER

R0

R1

P2

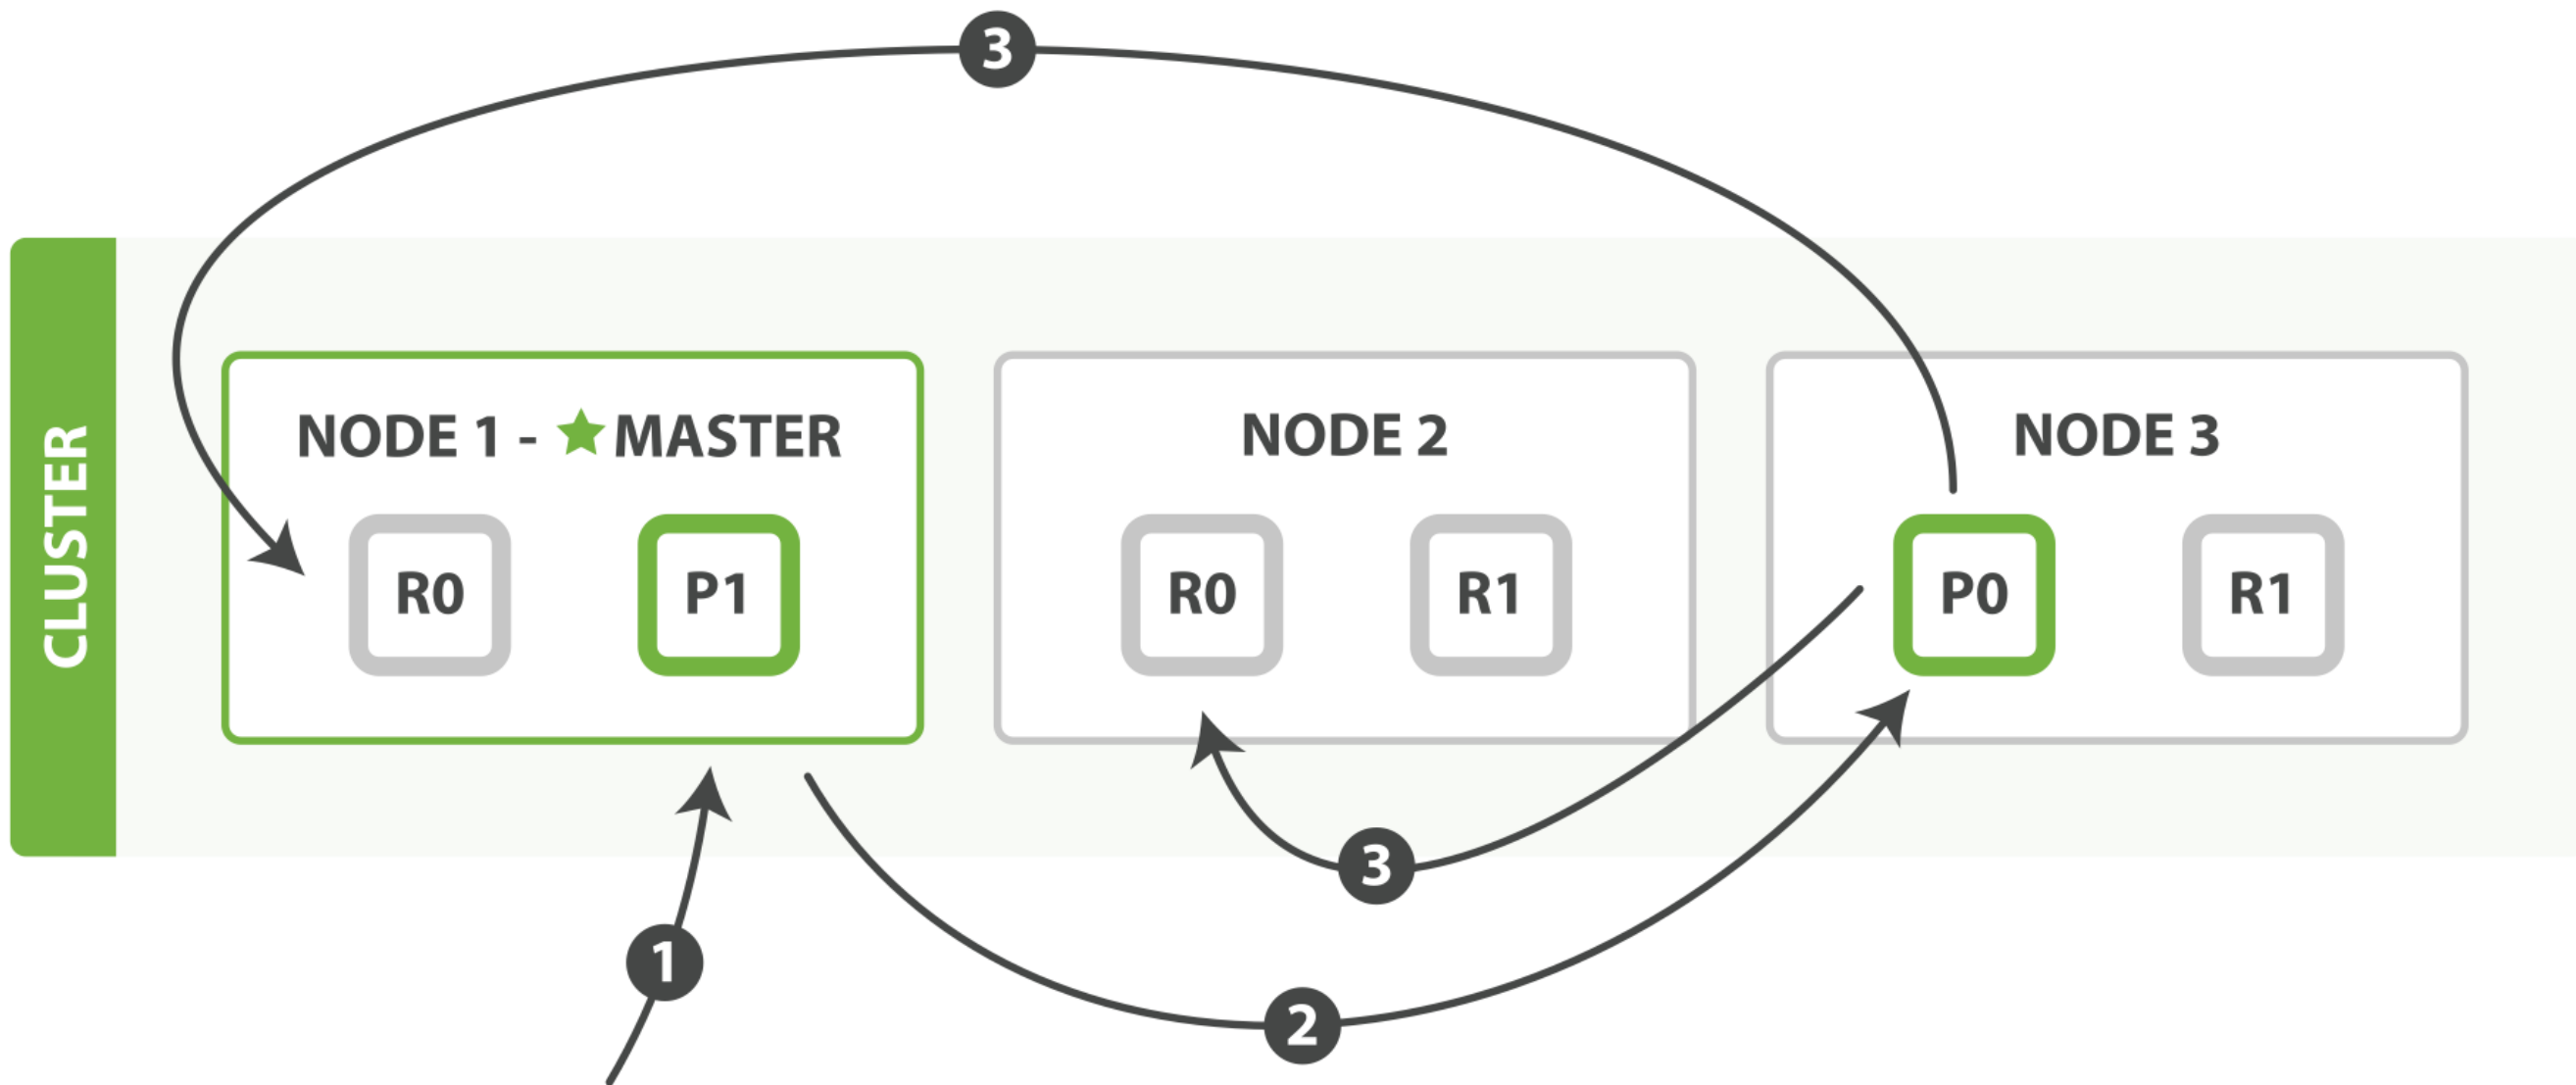
NODE 3

P0

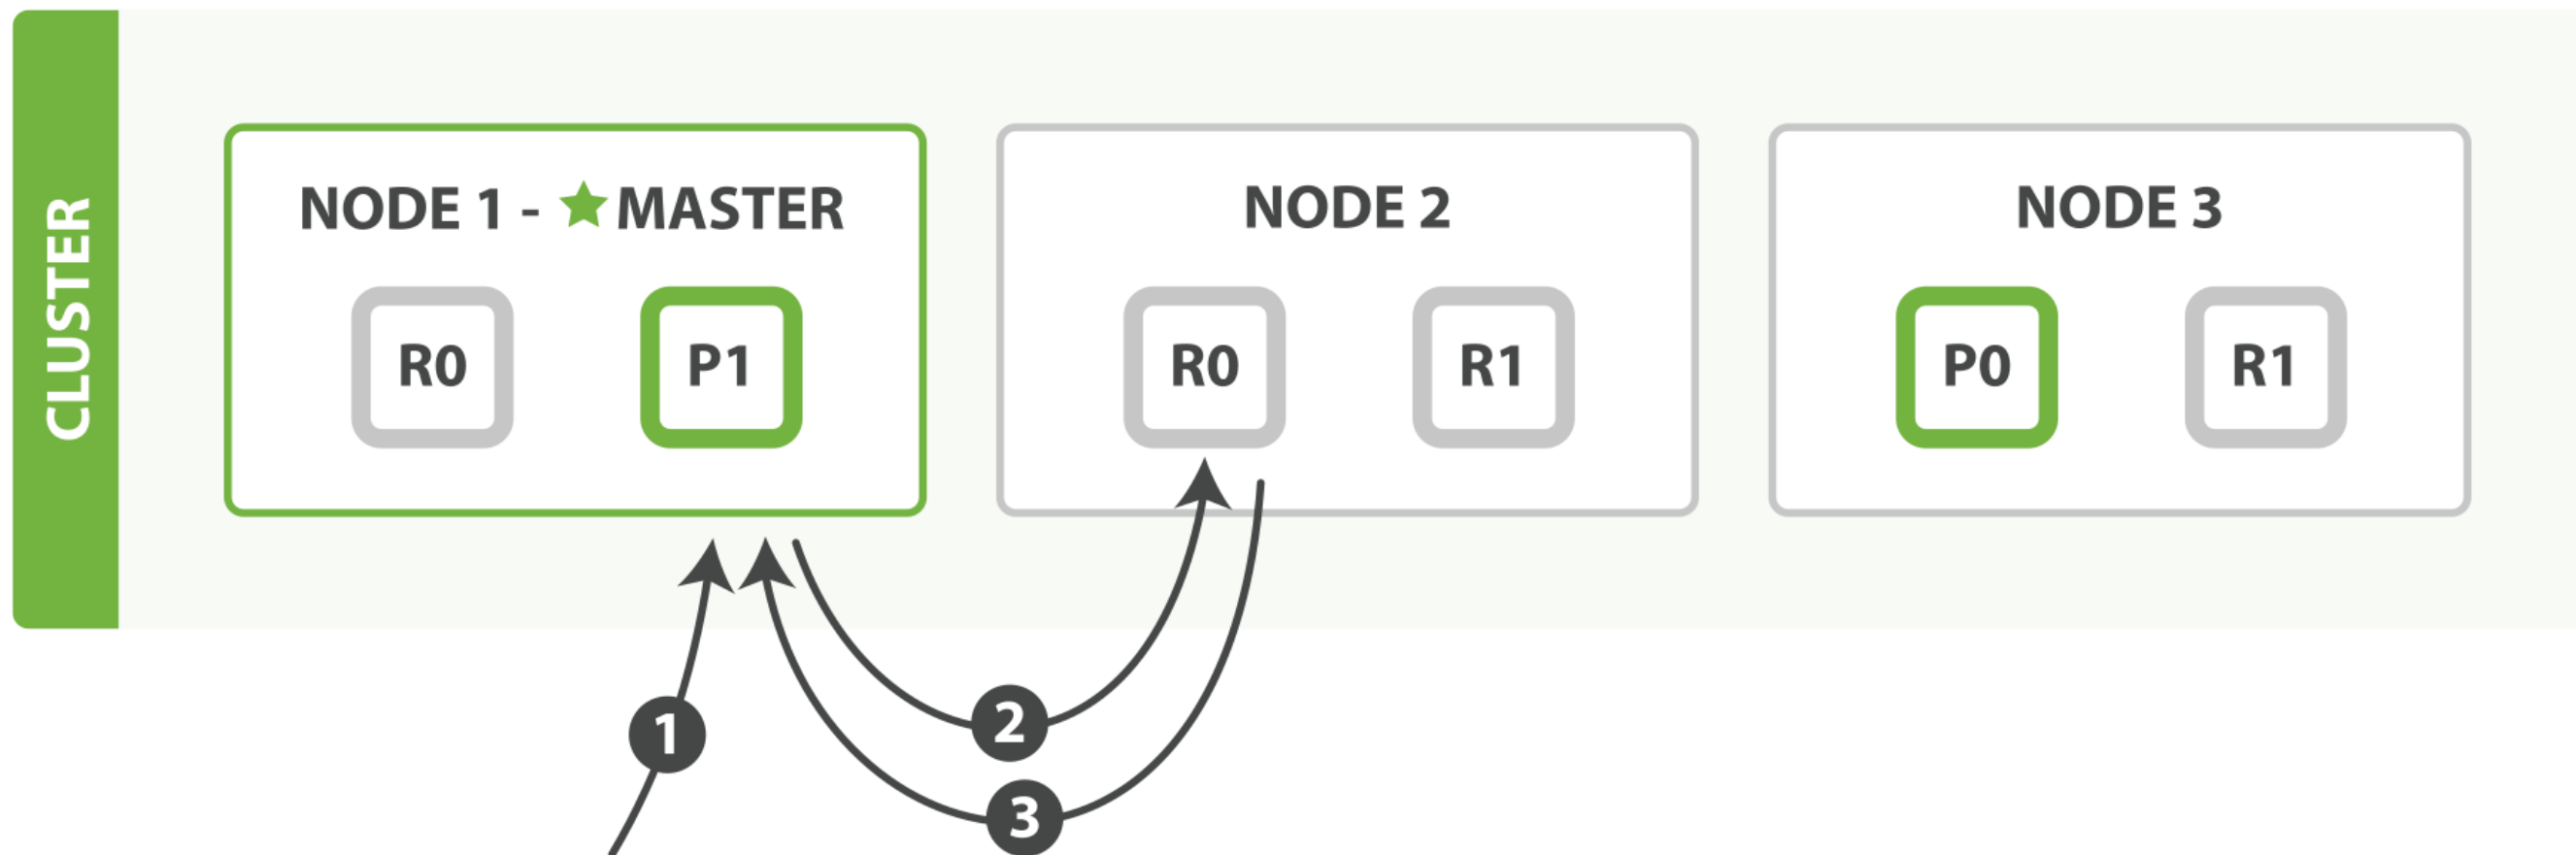
P1

R2

CREATING, INDEXING AND DELETING A DOCUMENT



RETRIEVING A DOCUMENT



4-2 RELATIONSHIPS

RELATIONSHIPS MATTER

- Blog Posts ↔ Comments
- Bank Accounts ↔ Transactions
- Orders ↔ Items
- Directories ↔ Files
- ...

SQL DATABASES

- ❑ Entities have an *unique primary key*
- ❑ *Normalization:*
 - ❑ Entity data is stored only once
 - ❑ Entities are referenced by primary key
- ❑ Updates happen in *only one place*

SQL DATABASES

□ Entities are *joined* at query time

```
SELECT Customer.name, Order.status  
FROM Order, Customer  
WHERE Order.customer_id = Customer.id
```

SQL DATABASES

- ❑ Changes are *ACID*
 - ❑ Atomicity
 - ❑ Consistency
 - ❑ Isolation
 - ❑ Durability

ATOMICITY

- ❑ If one part of the transaction fails, *the entire transaction fails*
- ❑ ...even in the event of power failure, crashes or errors
- ❑ *"all or nothing"*

CONSISTENCY

- ❑ Any transaction will bring the database from *one valid state to another*
- ❑ State must be valid according to all defined *rules*:
 - ❑ Constraints
 - ❑ Cascades
 - ❑ Triggers

ISOLATION

- ❑ The *concurrent execution* of transactions results in the same state that would be obtained if transactions were *executed serially*
- ❑ *Concurrency Control*

DURABILITY

- A transaction *will remain committed*
- ...even in the event of power failure, crashes or errors
- *Non-volatile memory*

SQL DATABASES

- ❑ Joining entities at query time is *expensive*
- ❑ *Impractical* with multiple nodes

- ❑ Treats the world as *flat*
- ❑ An index is a flat collection of *independent documents*
- ❑ A single document should contain all information to *match a search request*

- ❑ ACID support for changes on *single documents*
- ❑ No ACID transactions on *multiple documents*

ELASTICSEARCH

- ❑ Indexing and searching are *fast* and *lock-free*
- ❑ Massive amounts of data can be spread across *multiple nodes*

□ But we need *relationships*!

- ❑ Application-side joins
- ❑ Data denormalization
- ❑ Nested objects
- ❑ Parent/child relationships

4-3 APPLICATION-SIDE JOINS

APPLICATION-SIDE JOINS

- ❑ Emulates a relational database
- ❑ Joins at *application level*
- ❑ (*index*, *type*, *id*) = *primary key*

EXAMPLE

PUT /example/user/1

```
{  
  "name": "John Smith",  
  "email": "john@smith.com",  
  "born": "1970-10-24"  
}
```

EXAMPLE

PUT /example/blogpost/2

```
{  
  "title": "Relationships",  
  "body": "It's complicated",  
  "user": 1  
}
```


EXAMPLE

- (`example`, `user`, `1`) = *primary key*
- Store only the *id*
- *Index* and *type* are hard-coded into the application logic

EXAMPLE

GET /example/blogpost/_search

```
"query": {  
  "filtered": {  
    "filter": {  
      "term": { "user": 1 }  
    }  
  }  
}
```

EXAMPLE

- ❑ Blogposts written by “John”:
 - ❑ *Find ids* of users with name “John”
 - ❑ Find blogposts that *match the user ids*

EXAMPLE

GET /example/user/_search

```
"query": {  
  "match": {  
    "name": "John"  
  }  
}
```

EXAMPLE

□ For each *user id* from the first query:

GET /example/blogpost/_search

```
"query": {  
  "filtered": {  
    "filter": {  
      "term": { "user": <ID> }  
    }  
  }  
}
```

ADVANTAGES

- ❑ Data is *normalized*
- ❑ Change user data *in just one place*

DISADVANTAGES

- ❑ Run *extra queries* to join documents
- ❑ We could have *millions* of users named “John”
- ❑ *Less efficient* than SQL joins:
 - ❑ Several API requests
 - ❑ Harder to optimize

WHEN TO USE

- ❑ First entity has a *small number* of documents and they *hardly change*
- ❑ First query results can be *cached*

~~4-4~~ DATA DENORMALIZATION

DATA DENORMALIZATION

- ❑ *No joins*
- ❑ Store *redundant copies* of the data you need to query

EXAMPLE

PUT /example/user/1

```
{  
  "name": "John Smith",  
  "email": "john@smith.com",  
  "born": "1970-10-24"  
}
```

EXAMPLE

PUT /example/blogpost/2

```
{
  "title": "Relationships",
  "body": "It's complicated",
  "user": {
    "id": 1,
    "name": "John Smith"
  }
}
```

EXAMPLE

GET /example/blogpost/_search

```
"query": {  
  "bool": {  
    "must": [  
      { "match": {  
        "title": "relationships" } } ,  
      { "match": {  
        "user.name": "John" } }  
    ]  
  }  
}
```

ADVANTAGES

- *Speed*

- No need for expensive joins

DISADVANTAGES

- ❑ Uses more disk space (cheap)
- ❑ Update the same data in *several places*
 - ❑ *scroll* and *bulk* APIs can help
- ❑ *Concurrency issues*
 - ❑ *Locking* can help

WHEN TO USE

- ❑ Need for *fast search*
- ❑ Denormalized data *does not change* very often

4-5 NESTED OBJECTS

MOTIVATION

- ❑ Elasticsearch supports *ACID* when updating *single documents*
- ❑ Querying related data in the same document is *faster* (no joins)
- ❑ We want to *avoid denormalization*

THE PROBLEM WITH MULTILEVEL OBJECTS

PUT /example/blogpost/1

```
{  
  "title": "Nest eggs",  
  "body": "Making money...",  
  "tags": [ "cash", "shares" ],  
  "comments": [...]  
}
```

THE PROBLEM WITH MULTILEVEL OBJECTS

```
[{  
  "name": "John Smith",  
  "comment": "Great article",  
  "age": 28, "stars": 4,  
  "date": "2014-09-01"  
}, {  
  "name": "Alice White",  
  "comment": "More like this",  
  "age": 31, "stars": 5,  
  "date": "2014-10-22"  
}]
```

THE PROBLEM WITH MULTILEVEL OBJECTS

GET /example/blogpost/_search

```
"query": {  
  "bool": {  
    "must": [  
      {"match": {"name": "Alice"}},  
      {"match": {"age": "28"}}  
    ]  
  }  
}
```

THE PROBLEM WITH MULTILEVEL OBJECTS

```
[{  
  "name": "John Smith",  
  "comment": "Great article",  
  "age": 28, "stars": 4,  
  "date": "2014-09-01"  
}, {  
  "name": "Alice White",  
  "comment": "More like this",  
  "age": 31, "stars": 5,  
  "date": "2014-10-22"  
}]
```

THE PROBLEM WITH MULTILEVEL OBJECTS

- ❑ Alice is 31, not 28!
- ❑ It matched the age of John
- ❑ This is because indexed documents are stored as a *flattened dictionary*
- ❑ The correlation between *Alice* and 31 is *irretrievably lost*

THE PROBLEM WITH MULTILEVEL OBJECTS

```
{ "title": [eggs, nest],  
  "body": [making, money],  
  "tags": [cash, shares],  
  "comments.name":  
    [alice, john, smith, white],  
  "comments.comment":  
    [article, great, like, more, this],  
  "comments.age": [28, 31],  
  "comments.stars": [4, 5],  
  "comments.date":  
    [2014-09-01, 2014-10-22]}
```


NESTED OBJECTS

- ❑ Nested objects are indexed as *hidden separate documents*
- ❑ Relationships are *preserved*
- ❑ Joining nested documents is *very fast*

NESTED OBJECTS

```
{ "comments.name": [john, smith],  
  "comments.comment": [article, great],  
  "comments.age": [28],  
  "comments.stars": [4],  
  "comments.date": [2014-09-01]}  
  
{ "comments.name": [alice, white],  
  "comments.comment": [like, more, this],  
  "comments.age": [31],  
  "comments.stars": [5],  
  "comments.date": [2014-10-22]}
```

NESTED OBJECTS

```
{  
  "title": [eggs, nest],  
  "body": [making, money],  
  "tags": [cash, shares]  
}
```

NESTED OBJECTS

- ❑ Need to be enabled by *updating the mapping* of the index

MAPPING A NESTED OBJECT

PUT /example

```
"mappings": {  
  "blogpost": { "properties": {  
    "comments": { "type": "nested",  
      "properties": {  
        "name": {"type": "string"},  
        "comment": {"type": "string"},  
        "age": {"type": "short"},  
        "stars": {"type": "short"},  
        "date": {"type": "date"}  
      }  
    }  
  }  
}
```

QUERYING A NESTED OBJECT

GET /example/blogpost/_search

```
"query": {  
  "bool": {  
    "must": [  
      {"match": {"title": "eggs"}}  
      {"nested": <NESTED QUERY>}  
    ]  
  }  
}
```

NESTED QUERY

```
"nested": {  
  "path": "comments",  
  "query": {  
    "bool": {  
      "must": [  
        {"match":  
          {"comments.name": "john"}},  
        {"match":  
          {"comments.age": 28}}  
      ]  
    }  
  }  
}
```

THERE'S MORE

- ❑ Nested *filters*
- ❑ Nested *aggregations*
- ❑ *Sorting* by nested fields

ADVANTAGES

- ❑ *Very fast* query-time joins
- ❑ *ACID support* (single documents)
- ❑ Convenient search using *nested queries*

DISADVANTAGES

- ❑ To add, change or delete a nested object, the *whole document* must be *reindexed*
- ❑ Search requests return the *whole document*

WHEN TO USE

- ❑ When there is one main entity with a *limited number of closely related entities*
 - ❑ Ex: blogposts and comments
- ❑ Inefficient if there are *too many* nested objects

4-6 PARENT-CHILD RELATIONSHIP

PARENT-CHILD RELATIONSHIP

- ❑ *One-to-many* relationship
- ❑ Similar to the nested model
- ❑ Nested objects live in the *same document*
- ❑ Parent and children are *completely separate documents*

EXAMPLE

- ❑ Company with *branches* and *employees*
- ❑ Branch is the *parent*
- ❑ Employee are *children*

EXAMPLE

PUT /company

```
"mappings": {  
  "branch": {},  
  "employee": {  
    "_parent": {  
      "type": "branch"  
    }  
  }  
}
```

EXAMPLE

PUT /company/branch/london

```
{  
  "name": "London Westminster",  
  "city": "London",  
  "country": "UK"  
}
```


EXAMPLE

PUT /company/employee/1?
parent=london

```
{  
  "name": "Alice Smith",  
  "born": "1970-10-24",  
  "hobby": "hiking"  
}
```

FINDING PARENTS BY THEIR CHILDREN

GET /company/branch/_search

```
"query": {  
  "has_child": {  
    "type": "employee",  
    "query": {  
      "range": {  
        "born": {  
          "gte": "1980-01-01" }  
        }  
      }  
    }  
  }  
}
```

FINDING CHILDREN BY THEIR PARENTS

GET /company/employee/_search

```
"query": {  
  "has_parent": {  
    "type": "branch",  
    "query": {  
      "match": {  
        "country": "UK" }  
    }  
  }  
}
```

THERE'S MORE

- min_children and max_children
- Children *aggregations*
- *Grandparents* and *grandchildren*

ADVANTAGES

- ❑ Parent document can be updated *without reindexing* the children
- ❑ Child documents can be updated *without affecting* the parent
- ❑ Child documents can be returned in *search results* without the parent

ADVANTAGES

- ❑ Parent and children live on the *same shard*
- ❑ *Faster* than application-side joins

DISADVANTAGES

- ❑ Parent document and *all of its children* must live on the *same shard*
- ❑ 5 to 10 times *slower than nested queries*


WHEN TO USE

- ❑ *One-to-many* relationships
- ❑ When *index-time* is more important than *search-time* performance
- ❑ Otherwise, use *nested objects*

REFERENCES

MAIN REFERENCE

- ❑ *Elasticsearch, The Definitive guide*
- ❑ *Gormley & Tong*
- ❑ *O'Reilly*



OTHER REFERENCES

- ❑ *"Jepsen: simulating network partitions in DBs"*, <http://github.com/aphyr/jepsen>
- ❑ *"Call me maybe: Elasticsearch 1.5.0"*, <http://aphyr.com/posts/323-call-me-maybe-elasticsearch-1-5-0>
- ❑ *"Call me maybe: MongoDB stale reads"*, <http://aphyr.com/posts/322-call-me-maybe-mongodb-stale-reads>

OTHER REFERENCES

- ❑ *"Elasticsearch Data Resiliency Status"*,
<http://www.elastic.co/guide/en/elasticsearch/resiliency/current/index.html>
- ❑ *"Solr vs. Elasticsearch — How to Decide?"*,
<http://blog.sematext.com/2015/01/30/solr-elasticsearch-comparison/>

OTHER REFERENCES

- ❑ *"Changing Mapping with Zero Downtime"*,
<http://www.elastic.co/blog/changing-mapping-with-zero-downtime>

THANK YOU

Felipe Dornelas
felipedornelas.com
@felipead

ThoughtWorks®